

DYNOTEARS: Structure Learning from Time-Series Data


Daniele Montagnani

21 January 2024

Introduction

The Problem: Structure Learning for Dynamic Bayesian Networks

- What is structure learning?
- Static vs Dynamic
- Score-based vs Constraint-based



SOTA

- Heuristics to enforce acyclicity
- Learning relationships of different types either independently or sequentially

Paper Contribution

- NOTEARS framework was introduced to deal with static settings
- Adapt it to deal with time series data
- Simultaneous learning of contemporaneous and time lagged dependencies

Formulation

M independent realisations of a stationary time series

Variables influenced by other variables at present time and by previous observations up to the **autoregressive order** p

Structural Equation Model:

$$\mathbf{X} = \mathbf{XW} + \mathbf{YA} + \mathbf{Z}$$

Observations: $\mathbf{X} = [\dots | x_{m,t}^T | \dots]^T$

Time lagged observations: $\mathbf{Y} = [Y_1 | \dots | Y_p]$

Inter Slice Adjacency Matrices: $\mathbf{A} = [A_1^T | \dots | A_p^T]^T$

Errors: \mathbf{Z} - respecting proper conditions for identifiability

let's say, standard Gaussian, i.e. $z_{m,t} \sim N(0, I)$

Optimisation I

$$\min_{W, A} f(W, A) = \frac{1}{2n} \|X - XW - YA\|_F^2 + \lambda_W \|W\|_1 + \lambda_A \|A\|_1 \quad \text{s.t. } W \text{ is acyclic}$$

- Squared Error Loss
- l_1 regularization on W and A

How to Enforce Acyclicity in W ?

From NOTEARS: require $h(W) = 0$, where $h(W) = \text{tr } e^{W \circ W} - d$

Obtain a measure with the following properties:

1. Null iff W is acyclic
2. Quantifies the “Dagness” of the graph
3. Smooth
4. Easy to compute, like its derivatives

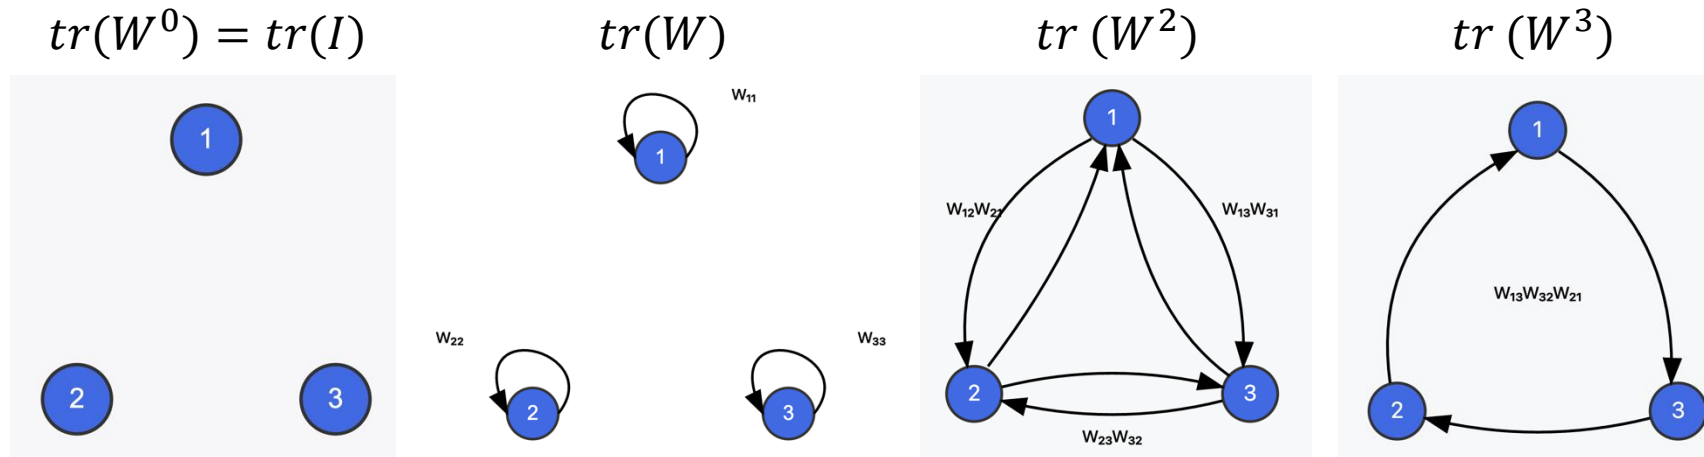
Yes, but why?

A beautiful math result

$$h(W) = \text{tr } e^{W \circ W} - d = 0 \quad \Rightarrow \quad \textit{Acyclic Graph}$$

But why?

- Why exponential? $e^x = \sum_{n=0}^{\infty} \frac{x^n}{n!} = 1 + x + \frac{x^2}{2} + \frac{x^3}{6} + \dots$
- Meaning of matrix multiplication and matrix powers
- Why trace?



- Why the Hadamard product? Isn't contraction a problem?

Optimisation II

After having introduced the constraint $h(W) = 0$

Solve with the Augmented Lagrange Method

$$\alpha^{k+1} = \alpha^k + \rho h(x^{k+1})$$

Obtain a series of smooth augmented objectives of the form

$$F(W, A) = f(W, A) + \frac{\rho}{2} h(W)^2 + \alpha h(W)$$

Equivalent to a quadratic program with non-negative constraints

For numerical stability:

eliminate weights close to 0 by thresholding the entries of W and A

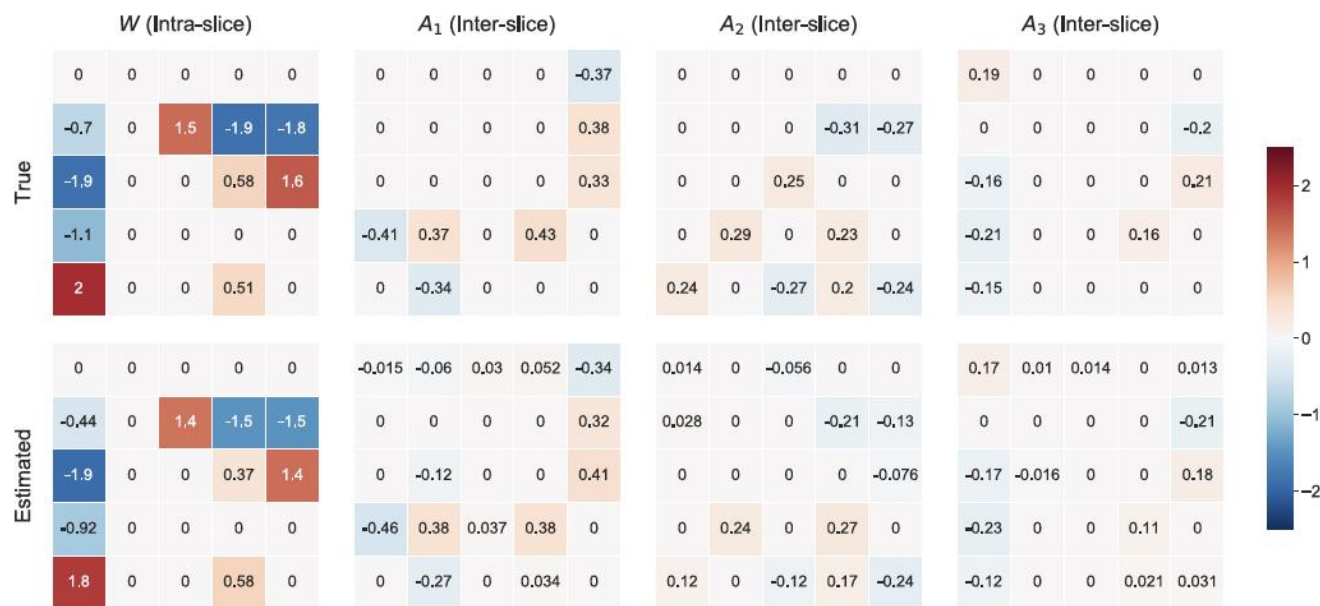
In the end, 4 total hyperparameters: $\lambda_W, \lambda_A, \tau_W, \tau_A$

Experiments on Synthetic Data

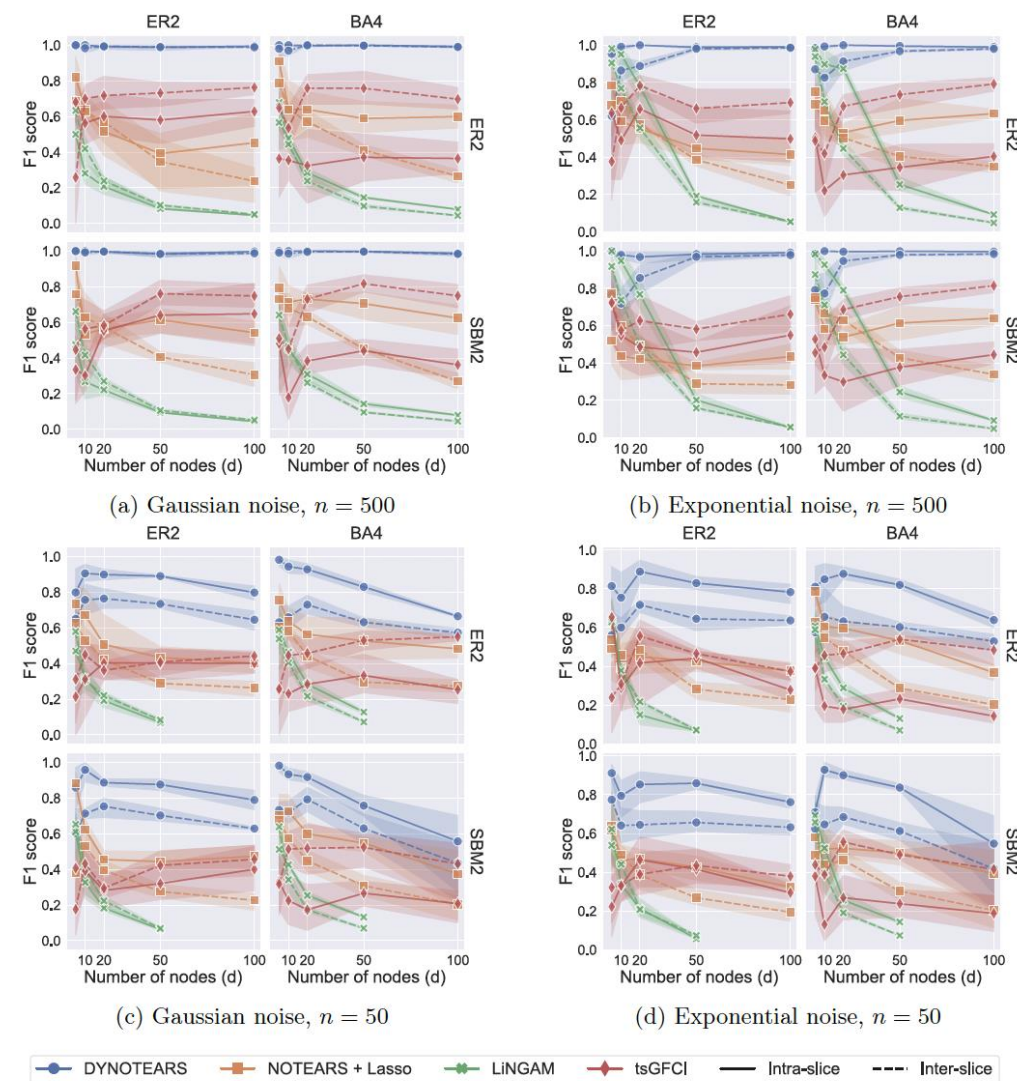
Simulation experiment:

1. Generate weighed graphs G_A and G_W
2. Generate synthetic data X and Y from these graphs
3. Run and eval DYNOTEARS on the synthetic data
4. Inspect the estimated weights
5. Compare with benchmarks from other algorithms

Estimated Weights (Gaussian noise, $n=500$, $d=5$, $p=3$)



Performances against benchmarking algorithms




Applications – S&P100 and DREAM4 Datasets

S&P100

- Use Normalised log returns
- Grid search → Optimal Hyperparameters
- No inter-slice weights
- Community detection algorithm → sector relationship

DREAM4

- Learn gene regulatory networks
- Use gene expression data
- Compare with flexible non-parametric methods
- Understand validity of the linear model assumption



Algorithm	Mean AUPR	Mean AUROC
DYNOTEARS	0.173	0.664
G1DBN	0.110	0.676
ScanBMA	0.101	0.657
VBSSMb	0.096	0.618
VBSSMa	0.086	0.624
Ebdbnet	0.043	0.643

Conclusion

The core of the paper was the application of the NOTEARS framework:

- Extremely elegant approach
- Obtains an optimization-friendly condition from a tricky to check requirement

The model rests on the fundamental assumption of linearity:

- Need to test its validity in different contexts
- Future directions can investigate non-linear approaches
- For example, with neural networks
- An approach reminiscent of Graph Neural Networks